



KING KHALID UNIVERSITY
COLLEGE OF COMPUTER SCIENCE

Detect Credit Card Fraud

By

Ahlam Mohammed Saad 441800053

Haifa Ibrahim al-mana 441813442

Supervised by:

MS.Mona Mushabab

Table of contents

1.1 Problem Statement	3
1.2 The Data Set.....	3
1.3 The Analysis	3
3.1 Importing the Datasets.....	3
3.2 Data Exploration	4
3.3 Data Manipulation	7
3.4 Data Modeling	7
3.5 Model Evaluation.....	8
1.5 Summary	8

1.1 Problem Statement

The aim of this R project is to build a classifier that can detect credit card fraudulent transactions. We will use a variety of machine learning algorithms that will be able to discern fraudulent from non-fraudulent one.

1.2 The Data Set

The datasets contains transactions made by credit cards in September 2013 by european cardholders.

This dataset presents transactions that occurred in two days, where we have 492 frauds out of 284,807 transactions. The dataset is highly unbalanced, the positive class (frauds) account for 0.172% of all transactions.

It contains only numerical input variables which are the result of a PCA transformation. Unfortunately, due to confidentiality issues, we cannot provide the original features and more background information about the data. Features V1, V2, ... V28 are the principal components obtained with PCA, the only features which have not been transformed with PCA are 'Time' and 'Amount'. Feature 'Time' contains the seconds elapsed between each transaction and the first transaction in the dataset. The feature 'Amount' is the transaction Amount, this feature can be used for example-dependant cost-sensitive learning. Feature 'Class' is the response variable and it takes value 1 in case of fraud and 0 otherwise.

1.3 The Analysis

3.1 Importing the Datasets

We are importing the datasets that contain transactions made by credit cards

The dataset used in this project is available here – [Fraud Detection Dataset](#)

3.2 Data Exploration

In this section of the fraud detection project, we explore the data that is contained in the creditcard_data dataframe. We will proceed by displaying the creditcard_data using the head() function. We will then proceed to explore the other components of this dataframe like summary(data\$Amount) it show us the "MIN, 1ST QU, MEDINE , MEAN, 3RD QU, MAX".

```
head(data)
# A tibble: 6 x 31
  Time    V1      V2      V3      V4      V5      V6      V7      V8      V9      V10     V11     V12
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 0 -1.36 -0.0728 2.54 1.38 -0.338 0.462 0.240 0.0987 0.364 0.0908 -0.552 -0.618
2 0 1.19 0.266 0.166 0.448 0.0600 -0.0824 -0.0788 0.0851 -0.255 -0.167 1.61 1.07
3 1 -1.36 -1.34 1.77 0.380 -0.503 1.80 0.791 0.248 -1.51 0.208 0.625 0.0661
4 1 -0.966 -0.185 1.79 -0.863 -0.0103 1.25 0.238 0.377 -1.39 -0.0550 -0.226 0.178
5 2 -1.16 0.878 1.55 0.403 -0.407 0.0959 0.593 -0.271 0.818 0.753 -0.823 0.538
6 2 -0.426 0.961 1.14 -0.168 0.421 -0.0297 0.476 0.260 -0.569 -0.371 1.34 0.360
... with 18 more variables: V13 <dbl>, V14 <dbl>, V15 <dbl>, V16 <dbl>, V17 <dbl>, V18 <dbl>,
V19 <dbl>, V20 <dbl>, V21 <dbl>, V22 <dbl>, V23 <dbl>, V24 <dbl>, V25 <dbl>, V26 <dbl>,
V27 <dbl>, V28 <dbl>, Amount <dbl>, Class <dbl>
apply(data, 2, function(x) sum(is.na(x)))
# A tibble: 31 x 1
  V1 V2 V3 V4 V5 V6 V7 V8 V9 V10 V11 V12 V13 V14 V15 V16 V17 V18 V19 V20 V21 V22 V23 V24 V25 V26 V27
  <dbl>
1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
2 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
3 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
4 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
5 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
6 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
7 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
8 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
9 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
10 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
11 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
12 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
13 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
14 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
15 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
16 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
17 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
18 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
19 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
20 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
21 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
22 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
23 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
24 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
25 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
26 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
27 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
28 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
29 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
30 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
31 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
```

Figure 1

As shown in figure 1 All the features, apart from "time" and "amount" are anonymised. Let's see whether there is any missing data. and There are no NA values in the data.

```
# summary(data$Amount)
# A tibble: 1 x 6
  Min. 1st Qu. Median Mean 3rd Qu. Max.
  <dbl> <dbl> <dbl> <dbl> <dbl> <dbl>
1 0.00 5.60 22.00 88.35 77.17 25691.16
```

Figure 2

As shown in figure 2 we can see the different statistics for our dataset

Visualization for the data set

By visualizing the data, we tried to find the range of amounts for the balances that were defrauded in our data set, and we found that the amounts range from (20,000 to 60,000)

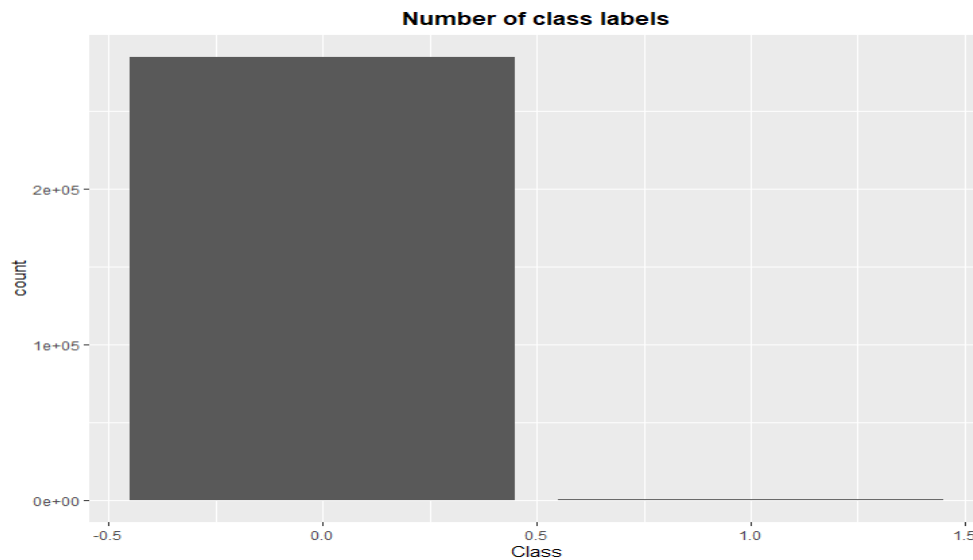


Figure 3

As shown in figure 3, clearly, the dataset is extremely unbalanced. Even a “null” classifier which always predicts class=0 would obtain over 99% accuracy on this task.

```
> print(p)
> summary(data)
```

Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	V10	V11	V12	V13	V14	V15	V16	V17	V18	V19	V20	V21	V22	V23
Min. : 0	Min. : -56.40751	Min. : -72.71573	Min. : -48.3256	Min. : -5.68317	Min. : -113.74331	Min. : -26.1605	Min. : -43.5572	Min. : -73.21672	Min. : -13.43407	Min. : -24.58826	Min. : -4.79747	Min. : -18.6837	Min. : -5.79188	Min. : -19.2143	Min. : -4.49894	Min. : -14.12985	Min. : -25.16280	Min. : -9.498746	Min. : -7.213527				
1st Qu.: 54202	1st Qu.: -0.92037	1st Qu.: -0.59855	1st Qu.: -0.8904	1st Qu.: -0.84864	1st Qu.: -0.69160	1st Qu.: -0.7683	1st Qu.: -0.5541	1st Qu.: -0.20863	1st Qu.: -0.64310	1st Qu.: -0.53543	1st Qu.: -0.76249	1st Qu.: -0.4056	1st Qu.: -0.64854	1st Qu.: -0.4256	1st Qu.: -0.58288	1st Qu.: -0.46804	1st Qu.: -0.48375	1st Qu.: -0.498850	1st Qu.: -0.456299				
Median : 84692	Median : 0.01811	Median : 0.06549	Median : 0.1799	Median : -0.01985	Median : -0.05434	Median : -0.2742	Median : 0.0401	Median : 0.02236	Median : -0.05143	Median : -0.09292	Median : -0.03276	Median : 0.1400	Median : -0.01357	Median : 0.0506	Median : 0.04807	Median : 0.06641	Median : -0.06568	Median : -0.003636	Median : 0.003735				
Mean : 94814	Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.0000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000	Mean : 0.0000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000	Mean : 0.00000				
3rd Qu.: 139321	3rd Qu.: 1.31564	3rd Qu.: 0.80372	3rd Qu.: 1.0272	3rd Qu.: 0.74334	3rd Qu.: 0.61193	3rd Qu.: 0.3986	3rd Qu.: 0.5704	3rd Qu.: 0.32735	3rd Qu.: 0.59714	3rd Qu.: 0.45392	3rd Qu.: 0.73959	3rd Qu.: 0.6182	3rd Qu.: 0.66251	3rd Qu.: 0.4931	3rd Qu.: 0.64882	3rd Qu.: 0.52330	3rd Qu.: 0.39968	3rd Qu.: 0.500807	3rd Qu.: 0.458949				
Max. : 172792	Max. : 2.45493	Max. : 22.05773	Max. : 9.3826	Max. : 16.87534	Max. : 34.80167	Max. : 73.3016	Max. : 120.5895	Max. : 20.00721	Max. : 15.59500	Max. : 23.74514	Max. : 12.01891	Max. : 7.8484	Max. : 7.12688	Max. : 10.5268	Max. : 8.87774	Max. : 17.31511	Max. : 9.25353	Max. : 5.041069	Max. : 5.591971				

Figure 4

As shown in figure 4, all the anonymised features seem to have been normalised with mean 0. We will apply that transformation to the “Amount” column.

Having normalized the “Amount” column, it is important to see how informative that feature would be in predicting whether a transaction was fraudulent. Hence, let’s plot the amount against the class of transaction.




Figure 5

There is clearly a lot more variability in the transaction values for non-fraudulent transactions. To get a fuller picture, let’s compute the mean and median values for each class. in figure 5

```
class 'mean(Amount)' 'median(Amount)'  
<dbl> <dbl> <dbl>  
1 0 88.3 22  
2 1 122. 9.25  
> |
```

Figure 6

fraudulent transactions seem to have higher mean value than non-fraudulent ones, meaning that this feature would likely be useful to use in the predictive model. However, the median is higher for the legitimate ones, meaning the distribution of values for class “0” is left-skewed (also seen on the figure 6).

3.3 Data Manipulation

In this section of the R data science project, we **normalize** our data. We will apply this to the amount component of data amount. **Normalization** is used to compensate for the differences in sample quantity (concentration and/or path length) between the database spectrum and the unknown spectrum

```
normalization
normalize <- function(x){
  return((x - mean(x, na.rm = TRUE))/sd(x, na.rm = TRUE))
}
data$Amount <- normalize(data$Amount)
--
```

Figure 7

3.4 Data Modeling

After we have normalized our entire dataset, we use logistic model. In this section of credit card fraud detection project, we will fit our model. A logistic regression is used for modeling the outcome probability of a class such as pass/fail, positive/negative and in our case – fraud/not fraud. We proceed to implement this model on our test data. We use Logistic Regression because it tries to minimize cost of how wrong a prediction is

```
#Logistic regression
log_mod <- glm(Class ~ ., family = "binomial", data = X_train)
summary(log_mod)
```

```
Cell:
glm(formula = Class ~ ., family = "binomial", data = X_train)

Deviance Residuals:
    Min       1Q   Median       3Q      Max
-4.8020  -0.0298  -0.0194  -0.0123   4.6025

Coefficients:
              Estimate Std. Error z value Pr(>|z|)
(Intercept) -8.274e+00  2.702e-01 -30.618 < 2e-16 ***
Time         -4.086e-06  2.485e-06  -1.645 0.100068
V1           9.452e-02  4.686e-02   2.017 0.043683 *
V2           1.804e-02  6.488e-02   0.278 0.780999
V3          -2.502e-02  5.889e-02  -0.425 0.671000
V4           6.845e-01  8.277e-02   8.271 < 2e-16 ***
V5           1.380e-01  7.473e-02   1.846 0.064821 .
V6          -1.070e-01  8.162e-02  -1.310 0.190035
V7          -9.144e-02  7.410e-02  -1.234 0.217190
V8          -1.699e-01  3.336e-02  -5.095 3.50e-07 ***
V9          -2.817e-01  1.249e-01  -2.255 0.024104 *
V10         -8.052e-01  1.054e-01  -7.640 2.16e-14 ***
V11         -1.338e-01  9.018e-02  -1.483 0.137993
V12          1.414e-01  1.013e-01   1.395 0.162906
V13         -3.470e-01  9.289e-02  -3.735 0.000188 ***
V14         -5.972e-01  7.105e-02  -8.406 < 2e-16 ***
V15         -6.793e-02  9.646e-02  -0.704 0.481283
```

Figure 8

As shown in figure 8 The model can be fitted using gradient descent on the parameter vector. Equipped with some basic information, let's see how the model performs.

3.5 Model Evaluation

We can evaluate a our model by doing the following:

The three main metrics used **to evaluate** a classification **model** are accuracy, precision, and recall.

```
# Use a threshold of 0.5 to transform predictions to binary
conf_mat <- confusionMatrix(y_test, as.numeric(predict(log_mod, X_test, type = "response") > 0.5))
print(conf_mat)
fourfoldplot(conf_mat$table)
```

Confusion Matrix and Statistics

	Reference	
Prediction	0	1
0	56861	8
1	34	58

Accuracy : 0.9993
95% CI : (0.999, 0.9995)
No Information Rate : 0.9988
P-Value [Acc > NIR] : 0.0010494

Kappa : 0.7338
McNemar's Test P-Value : 0.0001145

Sensitivity : 0.9994
Specificity : 0.8788
Pos Pred Value : 0.9999
Neg Pred Value : 0.6304
Prevalence : 0.9988
Detection Rate : 0.9982
Detection Prevalence : 0.9984
Balanced Accuracy : 0.9391

'Positive' Class : 0

Figure 9




Figure 10

As you can see in figure 10 that the end result of our project So the model predict 56861 true positive, 8 false positive, 34 false negatives, 58 true negatives.

Receiver Operating Characteristic (ROC) curve




Figure 11

As you can see in figure 11 the performance is very good with area under the curve is 0.984.

1.5 Summary

Concluding our R Data Science project, we learnt how to develop our credit card fraud detection model using machine learning. We used logistic model. We learnt how data can be analyzed and visualized to discern fraudulent transactions from other types of data.

1.6 Reference

1-<https://www.kaggle.com/mlg-ulb/creditcardfraud>

2- For the Source code on github click the link <https://github.com/Ahlam840/Credit-card>